

COMPUTER VISION AND PATTERN
RECOGNITION PROJECT
CNN CLASSIFIER

Alessandro Cesa

January 2024

Contents

Introduction	5
Problem statement	7
Implementation	9
Results	11
Bibliography	11

Introduction

The final project for the course in Computer Vision and Pattern Recognition consists in a classification task on a provided dataset containing 15 categories of images [2]. The classification has to be carried out by employing a Convolutional Neural Network, built following given specifications and steps described in the project assignment.

Problem statement

The project requires the implementation of an image classifier based on Convolutional Neural Networks, following three major steps.

1. Build from scratch a Convolutional Neural Network following this architecture:

Convolutional Neural Network Architecture

Layer	Size
Image Input	$64 \times 64 \times 1$
Convolution	$8 \times 3 \times 3$
ReLU	-
Max Pooling	2×2
Convolution	$16 \times 3 \times 3$
ReLU	-
Max Pooling	2×2
Convolution	$32 \times 3 \times 3$
ReLU	-
Fully Connected	15
Softmax	-
Classification Output	-

Following these specifications:

- Rescale the images to 64x64
- Split the provided training set in 85% for actual training set and 15% to be used as validation set;
- Employ the stochastic gradient descent with momentum optimization algorithm, using the default parameters of the library you use, except for those specified in the following;

- Use minibatches of size 32 and initial weights drawn from a Gaussian distribution with a mean of 0 and a standard deviation of 0.01; set the initial bias values to 0;

It's required to reach an accuracy of 30%

2. Improve the previous result, according to the following suggestions (not all the following will necessarily result in an increase of accuracy, but you should be able to obtain a test accuracy of about 60% by employing some of them):

- **Data Augmentation:** Given the small training set, data augmentation is likely to improve the performance. For the problem at hand, left-to-right reflections are a reasonable augmentation technique. By augmenting the training data using left-to-right reflections, you should get an accuracy of about 40%.
- **Batch Normalization [1]:** Add batch normalization layers before the ReLU layers.
- **Change Convolutional Filters:** Change the size and/or the number of the convolutional filters. For instance, try increasing their support as you move from input to output: 3×3 , 5×5 , 7×7 .
- **Optimization Parameters:** Play with the optimization parameters (learning rate, weights initialization, weights regularization, minibatch size, etc.). You can also switch to the Adam optimizer.
- **Dropout:** Add some dropout layers to improve regularization.
- **Ensemble of Networks:** Employ an ensemble of networks (five to ten), trained independently. Use the arithmetic average of the outputs to assign the class, as in [3].

- 3.

Implementation

Results

Bibliography

- [1] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International conference on machine learning*, pages 448–456. pmlr, 2015.
- [2] Svetlana Lazebnik, Cordelia Schmid, and Jean Ponce. Beyond bags of features: Spatial pyramid matching for recognizing natural scene categories. In *2006 IEEE computer society conference on computer vision and pattern recognition (CVPR'06)*, volume 2, pages 2169–2178. IEEE, 2006.
- [3] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1–9, 2015.